

Analysis

Sergio Rolando Valdivia Santisteban

March 2017

1 Introducción

Se compararan los resultados de algoritmos de multiplicación de matrices para ver cual de ellos presenta menos cache misses y tarda menos.

2 Comparación

```
sergio@sergio-Lenovo-Y50-70:~/Algoritmos-Paralelos/MultiplicacionMatrices$ valgrind --tool=cachegrind ./three
==4728== Cachegrind, a cache and branch-prediction profiler
==4728== Copyright (C) 2002-2015, and GNU GPL'd, by Nicholas Nethercote et al.
==4728== Using Valgrind-3.11.0 and LibVEX; rerun with -h for copyright info
==4728== Command: ./three
==4728==
--4728-- warning: L3 cache found, using its data for the LL simulation.
==4728==
==4728== I  refs:      2,158,679
==4728== I1 misses:    1,512
==4728== L1i misses:   1,445
==4728== I1 miss rate:  0.07%
==4728== L1i miss rate: 0.07%
==4728==
==4728== D  refs:      722,808 (531,428 rd + 191,380 wr)
==4728== D1 misses:   15,678 ( 13,487 rd +  2,191 wr)
==4728== L1d misses:   9,111 (  7,723 rd +  1,388 wr)
==4728== D1 miss rate:  2.2% (  2.5% +  1.1% )
==4728== L1d miss rate: 1.3% (  1.5% +  0.7% )
==4728==
==4728== LL refs:      17,190 ( 14,999 rd +  2,191 wr)
==4728== LL misses:   10,556 (  9,168 rd +  1,388 wr)
==4728== LL miss rate:  0.4% (  0.3% +  0.7% )
sergio@sergio-Lenovo-Y50-70:~/Algoritmos-Paralelos/MultiplicacionMatrices$ time ./three
real    0m0.001s
user    0m0.000s
sys     0m0.000s
sergio@sergio-Lenovo-Y50-70:~/Algoritmos-Paralelos/MultiplicacionMatrices$ █
```

Figure 1: Multipliacion simple de 3x3

```

sergio@sergio-Lenovo-Y50-70:~/Algoritmos-Paralelos/MultiplicacionMatrices$ valgrind --tool=cachegrind ./six
==4704== Cachegrind, a cache and branch-prediction profiler
==4704== Copyright (C) 2002-2015, and GNU GPL'd, by Nicholas Nethercote et al.
==4704== Using Valgrind-3.11.0 and LibVEX; rerun with -h for copyright info
==4704== Command: ./six
==4704==
--4704-- warning: L3 cache found, using its data for the LL simulation.
-----
==4704==
==4704== I   refs:      2,174,735
==4704== I1 misses:      1,642
==4704== L1i misses:     1,547
==4704== I1 miss rate:    0.08%
==4704== L1i miss rate:  0.07%
==4704==
==4704== D   refs:      728,805 (535,476 rd + 193,329 wr)
==4704== D1 misses:     15,830 ( 13,629 rd +  2,201 wr)
==4704== L1d misses:     9,152 (  7,763 rd +  1,389 wr)
==4704== D1 miss rate:    2.2% (  2.5% +  1.1% )
==4704== L1d miss rate:   1.3% (  1.4% +  0.7% )
==4704==
==4704== LL refs:        17,472 ( 15,271 rd +  2,201 wr)
==4704== LL misses:     10,699 (  9,310 rd +  1,389 wr)
==4704== LL miss rate:   0.4% (  0.3% +  0.7% )
sergio@sergio-Lenovo-Y50-70:~/Algoritmos-Paralelos/MultiplicacionMatrices$ time ./six
-----
real    0m0.002s
user    0m0.000s
sys     0m0.000s
sergio@sergio-Lenovo-Y50-70:~/Algoritmos-Paralelos/MultiplicacionMatrices$

```

Figure 2: Multipliacion por bloques de 3x3

Para matrices pequeñas el algoritmo de la multiplicacion simple termina en menos tiempo que el de la multiplicacion por bloques y parece ser que el primero tiene un mejor manejo del cache porque presenta menos cache misses. Al aumentar el tamaño del problema:

```

sergio@sergio-Lenovo-Y50-70:~/Algoritmos-Paralelos/MultiplicacionMatrices$ valgrind --tool=cachegrind ./three
==4867== Cachegrind, a cache and branch-prediction profiler
==4867== Copyright (C) 2002-2015, and GNU GPL'd, by Nicholas Nethercote et al.
==4867== Using Valgrind-3.11.0 and LibVEX; rerun with -h for copyright info
==4867== Command: ./three
==4867==
--4867-- warning: L3 cache found, using its data for the LL simulation.
-----
==4867==
==4867== I   refs:     53,974,798
==4867== I1 misses:       1,511
==4867== L1i misses:     1,444
==4867== I1 miss rate:    0.00%
==4867== L1i miss rate:  0.00%
==4867==
==4867== D   refs:     24,226,087 (22,874,975 rd + 1,351,112 wr)
==4867== D1 misses:      83,461 ( 79,670 rd +  3,791 wr)
==4867== L1d misses:    11,093 (  8,273 rd +  2,820 wr)
==4867== D1 miss rate:   0.3% (  0.3% +  0.3% )
==4867== L1d miss rate:  0.0% (  0.0% +  0.2% )
==4867==
==4867== LL refs:        84,972 ( 81,181 rd +  3,791 wr)
==4867== LL misses:     12,537 (  9,717 rd +  2,820 wr)
==4867== LL miss rate:   0.0% (  0.0% +  0.2% )
sergio@sergio-Lenovo-Y50-70:~/Algoritmos-Paralelos/MultiplicacionMatrices$ time ./three
-----
real    0m0.013s
user    0m0.012s
sys     0m0.000s
sergio@sergio-Lenovo-Y50-70:~/Algoritmos-Paralelos/MultiplicacionMatrices$

```

Figure 3: Multipliacion simple de 100x100

```
sergio@sergio-Lenovo-Y50-70:~/Algoritmos-Paralelos/MultiplicacionMatrices$ valgrind --tool=cachegrind ./six
==4845== Cachegrind, a cache and branch-prediction profiler
==4845== Copyright (C) 2002-2015, and GNU GPL'd, by Nicholas Nethercote et al.
==4845== Using Valgrind-3.11.0 and LibVEX; rerun with -h for copyright info
==4845== Command: ./six
==4845==
--4845-- warning: L3 cache found, using its data for the LL simulation.
-----
==4845==
==4845== I  refs:      55,021,188
==4845== I1 misses:    1,641
==4845== L1i misses:   1,546
==4845== I1 miss rate: 0.00%
==4845== L1i miss rate: 0.00%
==4845==
==4845== D  refs:      24,242,269 (22,889,111 rd + 1,353,158 wr)
==4845== D1 misses:    83,997 ( 80,170 rd + 3,827 wr)
==4845== L1d misses:   11,135 ( 8,313 rd + 2,822 wr)
==4845== D1 miss rate: 0.3% ( 0.4% + 0.3% )
==4845== L1d miss rate: 0.0% ( 0.0% + 0.2% )
==4845==
==4845== LL refs:       85,638 ( 81,811 rd + 3,827 wr)
==4845== LL misses:    12,681 ( 9,859 rd + 2,822 wr)
==4845== LL miss rate: 0.0% ( 0.0% + 0.2% )
sergio@sergio-Lenovo-Y50-70:~/Algoritmos-Paralelos/MultiplicacionMatrices$ time ./six
-----
real    0m0.012s
user    0m0.008s
sys     0m0.000s
sergio@sergio-Lenovo-Y50-70:~/Algoritmos-Paralelos/MultiplicacionMatrices$
```

Figure 4: Multipliacion simple de 100x100

Al aumentar el tamaño del problema el algoritmo de la multiplicacion simple resulta ser mas lento que el de la multiplicacion por bloques. a la vez el primero presenta un mayor numero de cahce misses mientras que el segundo menos, esto ocurre porque el algoritmo de la multiplicacion por bloques carga en cache porciones de memoria que lego sean usadas asiendo que el numero de cache misses sea mas bajo que el primero.

3 Conclusiones

- Tener en cuenta el manejo de la memoria cache puede ayudar a mejorar bastante el tiempo de un programa
- La multiplicación de matrices por bloques es mejor cuando se habla de varios elementos en la matris debido a que hace un mejor uso de la memoria cache.